



# Hands-on with MPI Programming and Spartan

Farzad Khodadadi  
PhD Candidate @ CLOUDS Lab  
The University of Melbourne  
[farzad.khodadadi@unimelb.edu.au](mailto:farzad.khodadadi@unimelb.edu.au)

# Outline

- MPI programming basics and common guidelines
- MPI4Py and MPJ Express demo
- Accessing Spartan, submitting jobs, and monitoring the results

# MPI Programming Basics

Many parallel programs can be written using just these six functions:

```
MPI_INIT  
MPI_FINALIZE  
MPI_COMM_SIZE  
MPI_COMM_RANK  
MPI_SEND  
MPI_RECV
```

- MPI\_SEND and MPI\_RECV functions can be substituted with collective operations such as MPI\_BCAST and MPI\_REDUCE

# Collective Operations in MPI

- ★ **MPI\_BCAST** distributes data from one process (the root) to all others in a communicator.
- ★ **MPI\_REDUCE** combines data from all processes in communicator and returns it to one process.
- ★ In many numerical algorithms, **SEND/RECEIVE** can be replaced by **BCAST/REDUCE**, improving both simplicity and efficiency.

# MPI4Py Sample Programs

```
from mpi4py import MPI  
import sys
```

```
size = MPI.COMM_WORLD.Get_size()  
rank = MPI.COMM_WORLD.Get_rank()  
print("Helloworld! I am process %d of %d.\n" % (rank,  
size))
```

```
import numpy as np
from mpi4py import MPI

from parutils import pprint

comm = MPI.COMM_WORLD

pprint("-"*78)
pprint(" Running on %d cores" % comm.size)
pprint("-"*78)

comm.Barrier()

# Prepare a vector of N=5 elements to be broadcasted...
N = 5
if comm.rank == 0:
    A = np.arange(N, dtype=np.float64)    # rank 0 has proper data
else:
    A = np.empty(N, dtype=np.float64)     # all other just an empty array

# Broadcast A from rank 0 to everybody
comm.Bcast( [A, MPI.DOUBLE] )

# Everybody should now have the same...
print "[%02d] %s" % (comm.rank, A)
```

---

```

import numpy as np
from mpi4py import MPI

from parutils import pprint

comm = MPI.COMM_WORLD

pprint("-"*78)
pprint(" Running on %d cores" % comm.size)
pprint("-"*78)

my_N = 4
N = my_N * comm.size

if comm.rank == 0:
    A = np.arange(N, dtype=np.float64)
else:
    A = np.empty(N, dtype=np.float64)

my_A = np.empty(my_N, dtype=np.float64)

# Scatter data into my_A arrays
comm.Scatter( [A, MPI.DOUBLE], [my_A, MPI.DOUBLE] )

pprint("After Scatter:")
for r in xrange(comm.size):
    if comm.rank == r:
        print "[%c] %s" % (comm.rank, my_A)
    comm.Barrier()

# Everybody is multiplying by 2
my_A *= 2

# Allgather data into A again
comm.Allgather( [my_A, MPI.DOUBLE], [A, MPI.DOUBLE] )

pprint("After Allgather:")
for r in xrange(comm.size):
    if comm.rank == r:
        print "[%c] %s" % (comm.rank, A)
    comm.Barrier()

```

# MPI Programming using MPJ Express

```
1 import mpi.*;
2
3 public class ToyExample {
4
5     public static void main(String[] args) throws Exception {
6
7         MPI.Init(args); int rank = MPI.COMM_WORLD.Rank(); int size = MPI.COMM_WORLD.Size() ;
8         int unitSize=4, tag=100, master=0;
9
10        if(rank == master) { /* master */
11
12            int sendbuf[] = new int[unitSize*(size-1)];
13
14            for(int i=1; i<size; i++)
15                MPI.COMM_WORLD.Send(sendbuf, (i-1)*unitSize, unitSize, MPI.INT, i, tag);
16
17            for(int i=1; i<size; i++)
18                MPI.COMM_WORLD.Recv(sendbuf, (i-1)*unitSize, unitSize, MPI.INT, i, tag);
19
20
21            for(int i=0 ; i<unitSize*(size-1) ; i++)
22                System.out.print(sendbuf[i]+" ");
23
24        } else { /* worker */
25
26            int recvbuf[] = new int[unitSize];
27            MPI.COMM_WORLD.Recv(recvbuf, 0, unitSize, MPI.INT, master, tag);
28
29            for(int i=0 ; i<unitSize; i++) recvbuf[i] = rank; /* computation loop */
30
31            MPI.COMM_WORLD.Send(recvbuf, 0, unitSize, MPI.INT, master, tag);
32        }
33
34        MPI.Finalize();
35    }
36 }
```



# MPI4Py and MPJ Express Demo



# Parallel Programming using Spartan

- Login to Spartan
  - `yourusername@spartan2.hpc.unimelb.edu.au`
- Upload your data (for your assignment, the data has been already uploaded and you only need to create a symbolic link to it in your home directory)
- Write a script to automate execution of your tasks
- Use SLURM's commands to submit your script, monitor your job's execution, cancel it, and much more.

# Spartan Demo



# More on SLURM and it's commands?

- <https://www.vlsci.org.au/documentation/managing-x86-jobs-slurm>
- <https://rc.fas.harvard.edu/resources/documentation/convenient-slurm-commands/>

A word cloud featuring the phrase "Thank You" in numerous languages. The words are arranged in a circular pattern, with "thank you" in the center in large blue letters. Other prominent words include "danke" (orange), "gracias" (red), "merci" (blue), "arigatō" (green), "sukriya" (green), "obrigado" (red), "dziękuje" (green), "bedankt" (blue), "tesekkür ederim" (green), "ngiyabonga" (blue), "tapadh leat" (blue), "moichchakkeram" (yellow), "go raibh maith agat" (green), "sagolun" (yellow), "chnorakaloutioun" (blue), "gratias ago" (blue), "gracies" (blue), "kop khun krap" (red), "taiku" (yellow), "arigatō" (green), "dakujem" (blue), "merci" (blue), "sukriya" (green), "kop khun krap" (red), "taiku" (yellow), "arigatō" (green), "dakujem" (blue), "merci" (blue), "sukriya" (green), "kop khun krap" (red), "taiku" (yellow), "arigatō" (green), "dakujem" (blue), "merci" (blue). The words are in various colors and sizes, creating a vibrant and multicultural visual.